

Java At Speed: Building a JVM For Modern Workloads

© Copyright Azul Systems 2015

Simon Ritter


Deputy CTO, Azul Systems

azul.com




@speakjava


Speed In The Java World



JVM Performance Graph: Ideal



JVM Performance Graph: Reality



Big JVM Challenges

Managed runtime environment

1. The Garbage Collector
 - Inherently non-deterministic
 - Pause times can be big for most algorithms
2. Bytecodes, not machine code
 - Adaptive compilation strategies
 - Speed of code 'warm-up'

What If There Was A Better JVM?



There is:
Zing

Azul Zing JVM

- Based on OpenJDK source code
- Passes all Java SE TCK/JCK tests
 - Drop in replacement for other JVMs
- Only one garbage collector: C4
 - Works in conjunction with Zing System Tools
- C2 JIT compiler replaced with Falcon
- ReadyNow! warm up elimination technology


Zing System Tools

- Enables better memory management for JVM
- Allocation of new pages comes from ZST
 - ZST knows cache status
 - Newly allocated pages for TLAB are 'hot'
 - Not like standard JVM
- Memory freed by JVM is returned to operating system
- Other clever tricks like Quick Release
- Only supported on Linux

Zing: Solving The GC Problem

- What Zing does NOT have
 - Serial, parallel, CMS or G1 collector
 - Full compacting old-generation fallback
 - Pause times proportional to heap size
- Zing: Continuous Concurrent Compacting Collector (C4)
 - Uses read-barrier rather than write-barrier
 - No stop-the-world pauses
 - The bigger the heap, the better the results

Azul Continuous Concurrent Compacting Collector (C4)



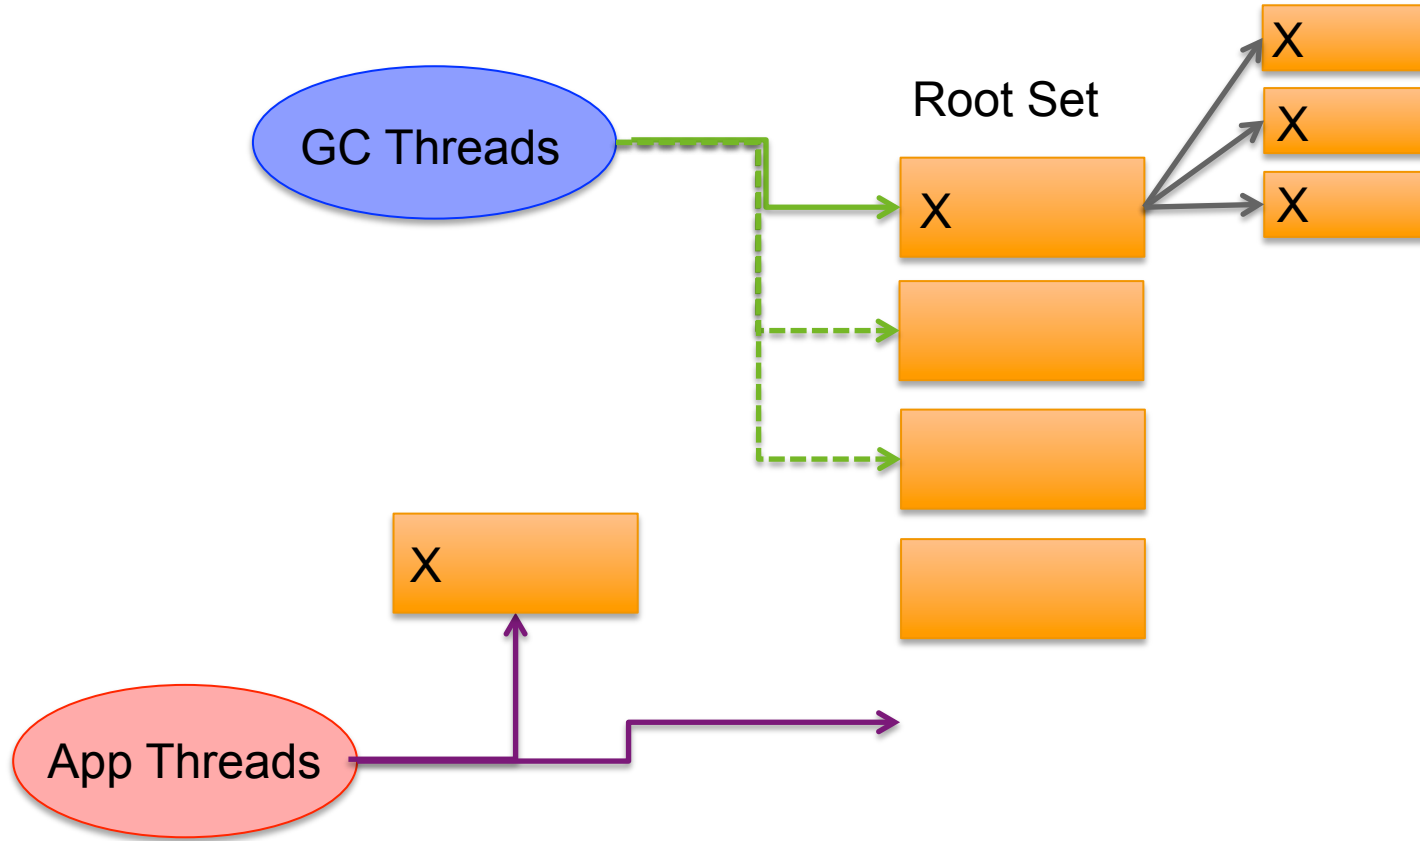
C4 Basics

- Generational (young and old)
 - Uses the same GC collector for both
 - For efficiency rather than pause containment
 - Weak generational hypothesis makes sense
- Concurrent, parallel and compacting
 - But no STW compacting fallback
- Algorithm is mark, relocate, remap


Loaded Value Barrier

- Read barrier
 - Tests all object references as they are loaded
- Enforces two invariants
 - Reference is marked through
 - Reference points to correct object position
- Allows for concurrent marking and relocation
- Minimal performance overhead
 - Test and jump (2 instructions)
 - x86 architecture reduces this to one micro-op

Concurrent Mark Phase




Loaded Value Barrier: Marking




```
if (phase == MARKING && B not marked) {  
    mark(B);  
    addGCRootSet(B);  
}  
  
return B;
```

Relocation Phase



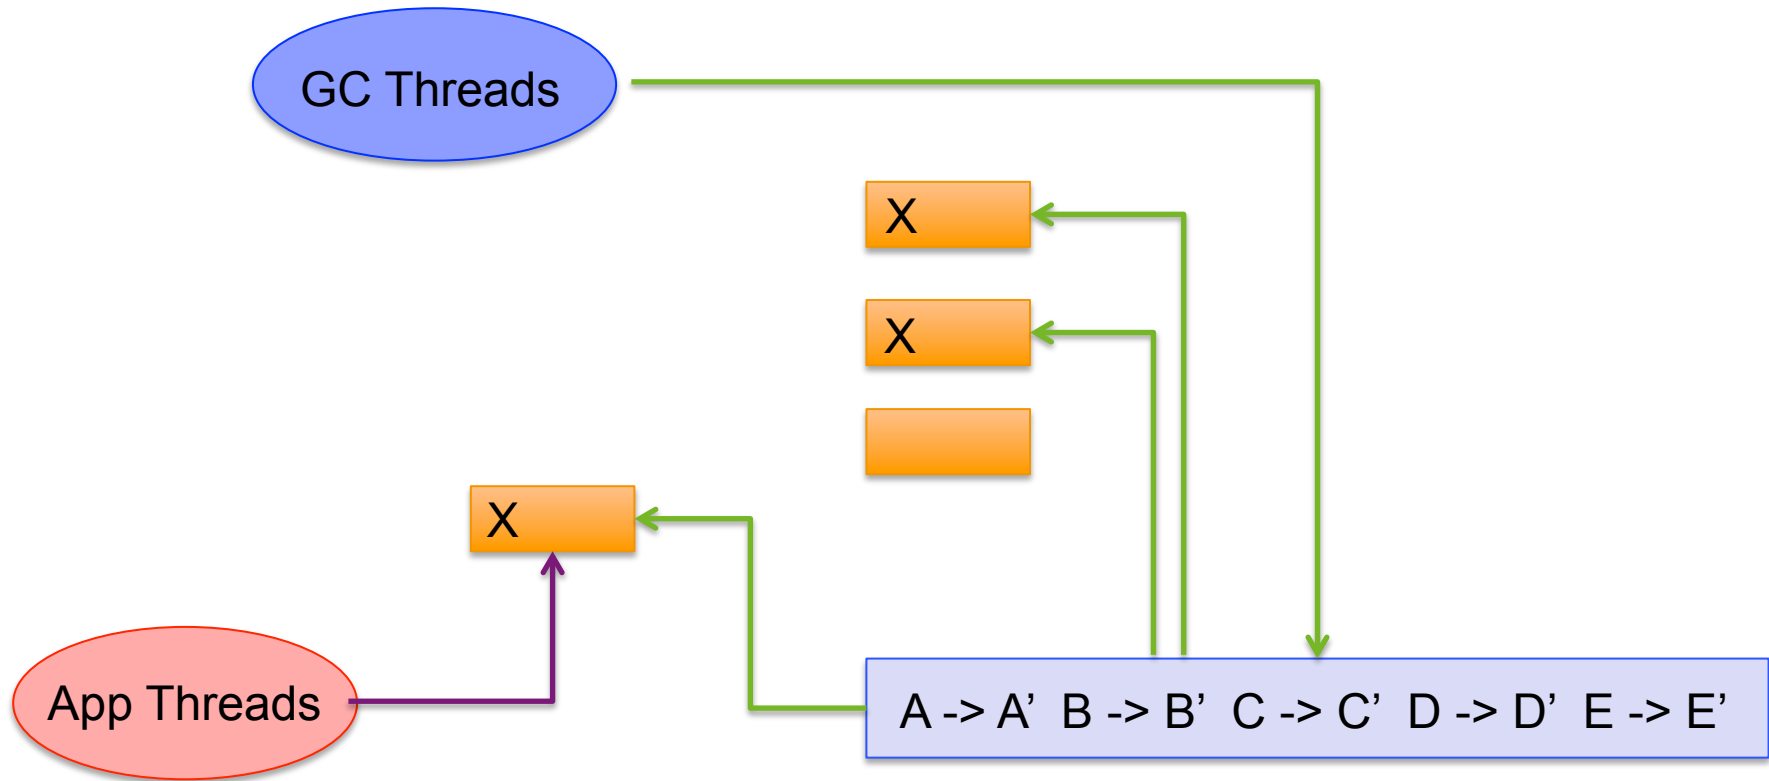
A -> A' B -> B' C -> C' D -> D' E -> E'

Quick Release




A -> A' B -> B' C -> C' D -> D' E -> E'

Remapping Phase



Loaded Value Barrier: Remapping




```
if (phase == REMAP && checkSideTable(B)) {  
    remap(B);  
    removeSideTable(B)  
}  
  
return B;
```

Zing: Big Heaps, No Problem

- Scales to 8Tb heap
 - No degradation in pause times
- Use one big heap, rather than many small heaps
 - Less JVMs means more efficiency
- Zing does not require big heaps
 - But works well with them

GC Tuning



Non-Zing GC Tuning Options



GC Tuning Used To Be Hard

```
Java -Xmx12g -XX:MaxPermSize=64M -XX:PermSize=32M -XX:MaxNewSize=2g  
-XX:NewSize=1g -XX:SurvivorRatio=128 -XX:+UseParNewGC  
-XX:+UseConcMarkSweepGC -XX:MaxTenuringThreshold=0  
-XX:CMSInitiatingOccupancyFraction=60 -XX:+CMSParallelRemarkEnabled  
-XX:+UseCMSInitiatingOccupancyOnly -XX:ParallelGCThreads=12  
-XX:LargePageSizeInBytes=256m ...
```

```
Java -Xms8g -Xmx8g -Xmn2g -XX:PermSize=64M -XX:MaxPermSize=256M  
-XX:-OmitStackTraceInFastThrow -XX:SurvivorRatio=2  
-XX:-UseAdaptiveSizePolicy -XX:+UseConcMarkSweepGC  
-XX:+CMSConcurrentMTEnabled -XX:+CMSParallelRemarkEnabled  
-XX:+CMSParallelSurvivorRemarkEnabled  
-XX:CMSMaxAbortablePrecleanTime=10000  
-XX:+UseCMSInitiatingOccupancyOnly  
-XX:CMSInitiatingOccupancyFraction=63 -XX:+UseParNewGC -Xnoclassgc ...
```

GC Tuning Used To Be Hard

```
Java -Xmx12g -XX:MaxPermSize=64M -XX:PermSize=32M -XX:MaxNewSize=2g  
-XX:NewSize=1g -XX:SurvivorRatio=128 -XX:+UseParNewGC  
-XX:+UseConcMarkSweepGC -XX:MaxTenuringThreshold=0  
-XX:CMSInitiatingOccupancyFraction=60 -XX:+CMSParallelRemarkEnabled  
-XX:+UseCMSInitiatingOccupancyOnly -XX:ParallelGCThreads=12  
-XX:LargePageSizeInBytes=256m ...
```

```
Java -Xms8g -Xmx8g -Xmn2g -XX:PermSize=64M -XX:MaxPermSize=256M  
-XX:-OmitStackTraceInFastThrow -XX:SurvivorRatio=2  
-XX:-UseAdaptiveSizePolicy -XX:+UseConcMarkSweepGC  
-XX:+CMSConcurrentMTEnabled -XX:+CMSParallelRemarkEnabled  
-XX:+CMSParallelSurvivorRemarkEnabled  
-XX:CMSMaxAbortablePrecleanTime=10000  
-XX:+UseCMSInitiatingOccupancyOnly  
-XX:CMSInitiatingOccupancyFraction=63 -XX:+UseParNewGC -Xnoclassgc ...
```

GC Tuning With Zing

java -Xmx1g

java -Xmx10g

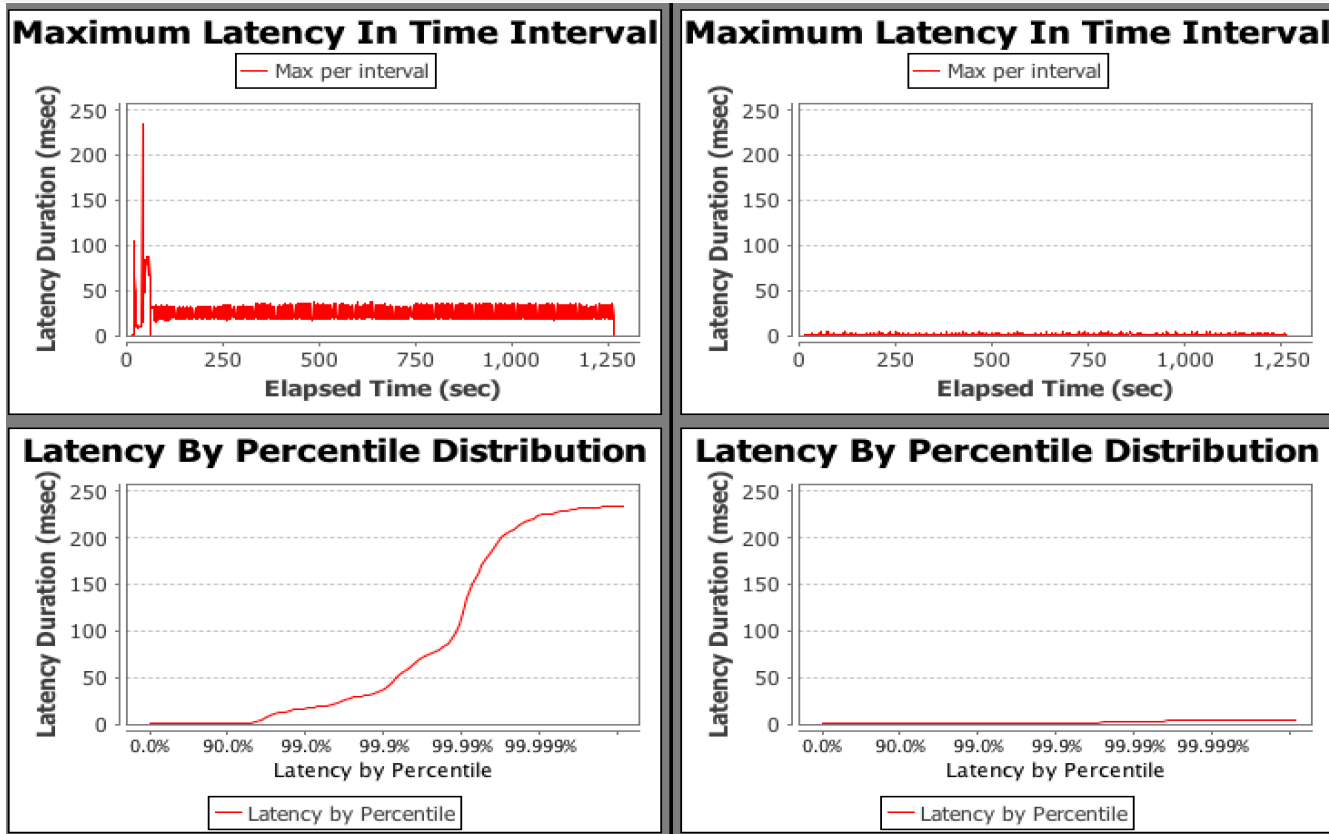
java -Xmx100g

java -Xmx8t

Measuring Platform Performance

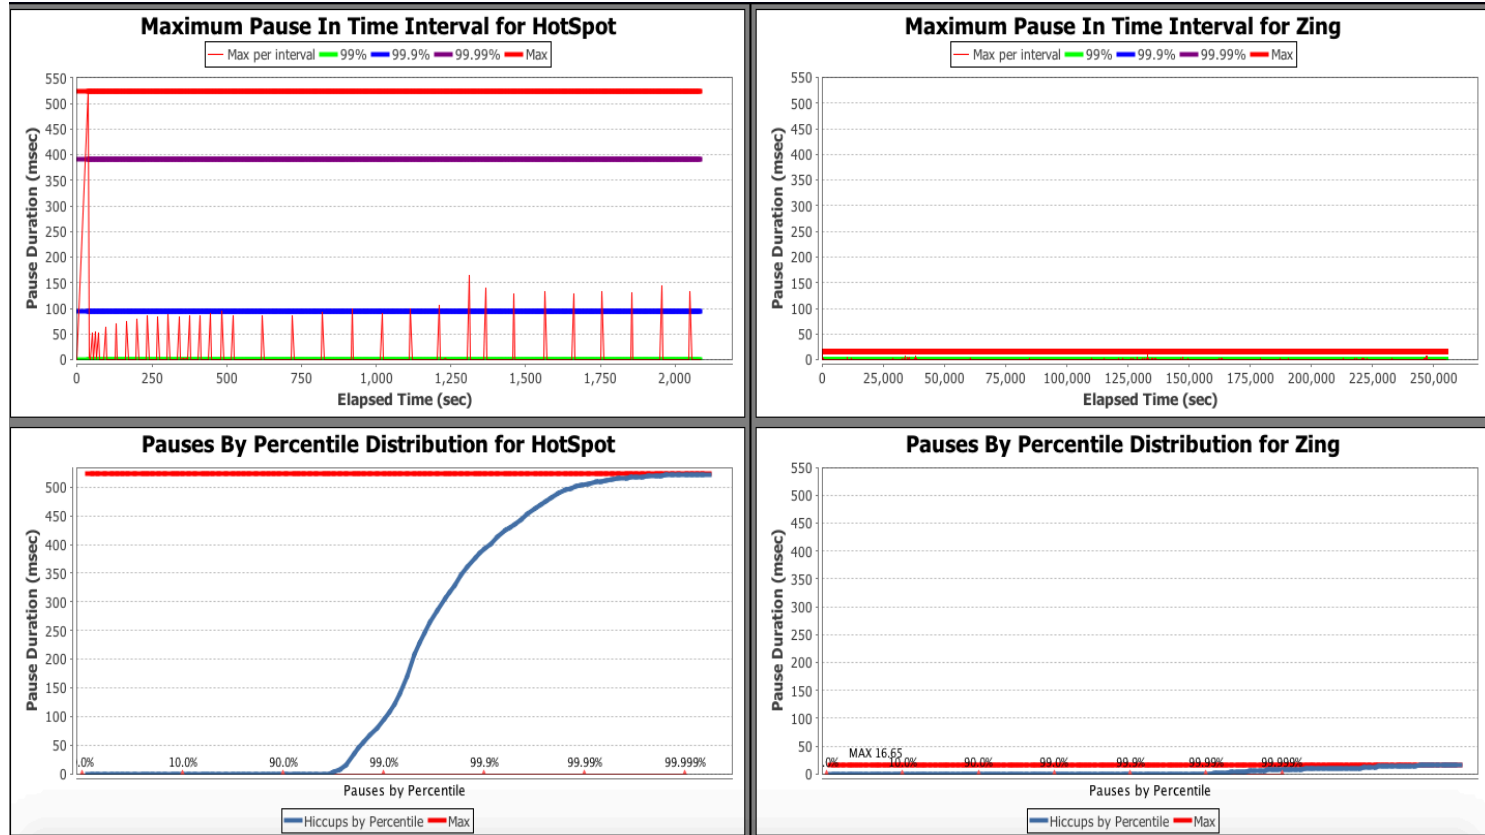
- jHiccup
 - Spends most of its time asleep
 - Minimal effect on performance
 - Wakes every 1 ms
 - Records delta of time it expects to wake up
 - Measured effect is what would be experienced by your application
 - Generates histogram log files
 - These can be graphed for easy evaluation

Small Heap, Small Latency




Hazelcast 2-node system with 1Gb heap Hotspot v. Zing

Big Heap, Small Latency



Cassandra with 60Gb heap Hotspot v. Zing

Azul Falcon JIT Compiler



AOT And JIT

- Ahead of Time (Static) compilation
 - Restricted optimisations due to runtime changes
 - Classloading
- Just in Time (Adaptive) compilation
 - Complete picture of loaded classes
 - More aggressive optimisations
 - Method inlining
 - Escape analysis

Speculative Optimisations

- JIT compilers can do things that static compilers are not able to do
 - However, these can sometimes affect performance
- Speculative optimisations
 - Optimise code based on what has happened so far
 - Likely to continue
 - But may not...
 - If speculation is false compiled code is thrown away
 - Revert to interpreted mode
 - Recompile if necessary

Untaken Path Example

```
int computeMagnitude(int val) {  
    if (val >= 10)  
        bias = computeBias(val);  
    else  
        bias = 1;  
    return Math.log10(bias + 99);  
}
```

Profiling shows val
always less than 10
(so far)



```
int computeMagnitude(int val) {  
    if (val >= 10)  
        uncommonTrap();  
  
    return 2;  
}
```

Speculative optimised
code compiled as if it
was this

Implicit Null Check

- All field and array references are implicitly null checked

```
x = foo.y;
```

is really compiled as

```
if (foo == null)
    throw new NullPointerException();
x = foo.y;
```

- Compiler can hope for non-nulls and handle the SEGV separately
 - Faster code path

Adaptive Compilation Challenges

- Traditionally three options for running bytecodes
 - Fully interpreted
 - C1 (Client compiler):
 - Fast warmup, lower optimal level
 - C2 (Server compiler):
 - Slower warmup, higher optimal level
- Application takes time from starting to optimal level of performance

Advancing Adaptive Compilation

- Azul Falcon JVM compiler
 - Based on latest compiler research
 - LLVM project
- Better performance
 - Better intrinsics
 - More inlining
 - Fewer compiler excludes
- Targeted as replacement for C2 compiler



Simple Code Example

- Simple array summing loop
 - A modern compiler will use vector operations for this

```
private int sumLoop(int[] a) {  
    int sum = 0;  
    for (int i = 0; i < a.length; i++) {  
        sum += a[i];  
    }  
    return sum;  
}
```

More Complex Code Example

- Conditional array cell addition loop
 - Hard for compiler to identify for vector instruction use

```
private void addArraysIfEven(int a[], int b[]) {  
    if (a.length != b.length) {  
        throw new RuntimeException("length mismatch");  
    }  
    for (int i = 0; i < a.length; i++) {  
        if ((b[i] & 0x1) == 0) {  
            a[i] += b[i];  
        }  
    }  
}
```

Traditional JVM JIT

Per element jumps
2 elements per iteration

```
private void addArraysIfEven(int a[], int b[]) {  
    if (a.length != b.length) {  
        throw new RuntimeException("length mismatch");  
    }  
    for (int i = 0; i < a.length; i++) {  
        if ((b[i] & 0x1) == 0) {  
            a[i] += b[i];  
        }  
    }  
}
```

		0x3001067f	addl %ecx, 12(%rsi)	0x014e0c
		0x30010682	movl \$1, %edi	0xbf01000000
		0x30010687	cmpl \$1, %eax	0x83f801
		0x3001068a	je 56 ; ABS: 0x300106c4	0x7438
		0x3001068c	subq %rdi, %rax	0x4829f8
		0x3001068f	leaq 16(%rdx,%rdi,4), %rcx	0x488d4cba10
		0x30010694	leaq 16(%rsi,%rdi,4), %rdx	0x488d54be10
		0x30010699	noop (%rax)	0x0f1f8000000000
16.84%	1,286	0x300106a0	movl -4(%rcx), %esi	0x8b71fc
6.10%	466	0x300106a3	testb \$1, %sil	0x40f6c601
		0x300106a7	jne 3 ; ABS: 0x300106ac	0x7503
7.61%	581	0x300106a9	addl %esi, -4(%rdx)	0x0172fc
29.41%	2,246	0x300106ac	movl (%rcx), %esi	0x8b31
2.25%	172	0x300106ae	testb \$1, %sil	0x40f6c601
		0x300106b2	jne 2 ; ABS: 0x300106b6	0x7502
8.00%	611	0x300106b4	addl %esi, (%rdx)	0x0132
29.73%	2,271	0x300106b6	addq \$8, %rcx	0x4883c108
		0x300106ba	addq \$8, %rdx	0x4883c208
		0x300106be	addq \$-2, %rax	0x4883c0fe
		0x300106c2	jne -36 ; ABS: 0x300106a0	0x75dc
0.03%	2	0x300106c4	addq \$24, %rsp	0x4883c418
0.03%	2	0x300106c8	retq	0xc3
		0x300106c9	movq %rsi, 16(%rsp)	0x4889742410
		0x300106ce	movq %rdx, 8(%rsp)	0x4889542408
		0x300106d3	movabsq \$805334400, %rax	0x48b8806d003000000000
		0x300106dd	callq *%rax	0xffd0

Falcon JIT

Using AVX2 vector instructions
32 elements per iteration



```
private void addArraysIfEven(int a[], int b[]) {  
    if (a.length != b.length) {  
        throw new RuntimeException("length mismatch");  
    }  
    for (int i = 0; i < a.length; i++) {  
        if ((b[i] & 0x1) == 0) {  
            a[i] += b[i];  
        }  
    }  
}
```


Broadwell E5-2690-v4

		0x3001455b	movq %rdi, %rbx
		0x3001455e	nop
0.15%	4	0x30014560	vmovdqu -96(%r11), %ymm2
12.31%	320	0x30014566	vmovdqu -64(%r11), %ymm3
0.50%	13	0x3001456c	vmovdqu -32(%r11), %ymm4
2.04%	53	0x30014572	vmovdqu (%r11), %ymm5
0.31%	8	0x30014577	vpand %ymm0, %ymm2, %ymm6
4.54%	118	0x3001457b	vpand %ymm0, %ymm3, %ymm7
0.69%	18	0x3001457f	vpand %ymm0, %ymm4, %ymm8
1.35%	35	0x30014583	vpand %ymm0, %ymm5, %ymm9
0.42%	11	0x30014587	vpcmpeqd %ymm1, %ymm6, %ymm6
2.58%	67	0x3001458b	vpmaskmovd -96(%rcx), %ymm6, %ymm10
3.58%	93	0x30014591	vpcmpeqd %ymm1, %ymm7, %ymm7
2.12%	55	0x30014595	vpmaskmovd -64(%rcx), %ymm7, %ymm11
12.12%	315	0x3001459b	vpcmpeqd %ymm1, %ymm8, %ymm8
1.50%	39	0x3001459f	vpmaskmovd -32(%rcx), %ymm8, %ymm12
3.69%	96	0x300145a5	vpcmpeqd %ymm1, %ymm9, %ymm9
1.81%	47	0x300145a9	vpmaskmovd (%rcx), %ymm9, %ymm13
12.27%	319	0x300145ae	vpadd %ymm2, %ymm10, %ymm2
0.58%	15	0x300145b2	vpadd %ymm3, %ymm11, %ymm3
0.19%	5	0x300145b6	vpadd %ymm4, %ymm12, %ymm4
0.58%	15	0x300145ba	vpadd %ymm5, %ymm13, %ymm5
3.27%	85	0x300145be	vpmaskmovd %ymm2, %ymm6, -96(%rcx)
7.15%	186	0x300145c4	vpmaskmovd %ymm3, %ymm7, -64(%rcx)
13.65%	355	0x300145ca	vpmaskmovd %ymm4, %ymm8, -32(%rcx)
4.58%	119	0x300145d0	vpmaskmovd %ymm5, %ymm9, (%rcx)
6.81%	177	0x300145d5	subq \$-128, %r11
0.69%	18	0x300145d9	subq \$-128, %rcx
0.31%	8	0x300145dd	addq \$-32, %rbx
		0x300145e1	jne -135 ; ABS: 0x30014560
		0x300145e7	testl %r9d, %r9d
		0x300145ea	jne -356 ; ABS: 0x3001448c

ReadyNow!




Traditional JVM



ReadyNow! Solution

- Save JVM JIT profiling information
 - Classes loaded
 - Classes initialised
 - Instruction profiling data
 - Speculative optimisation failure data
- Data can be gathered over much longer period
 - JVM/JIT profiles quickly
 - Significant reduction in deoptimisations
- Able to load, initialise and compile most code before `main()`

Effect Of ReadyNow!



Customer application

Compile Stashing

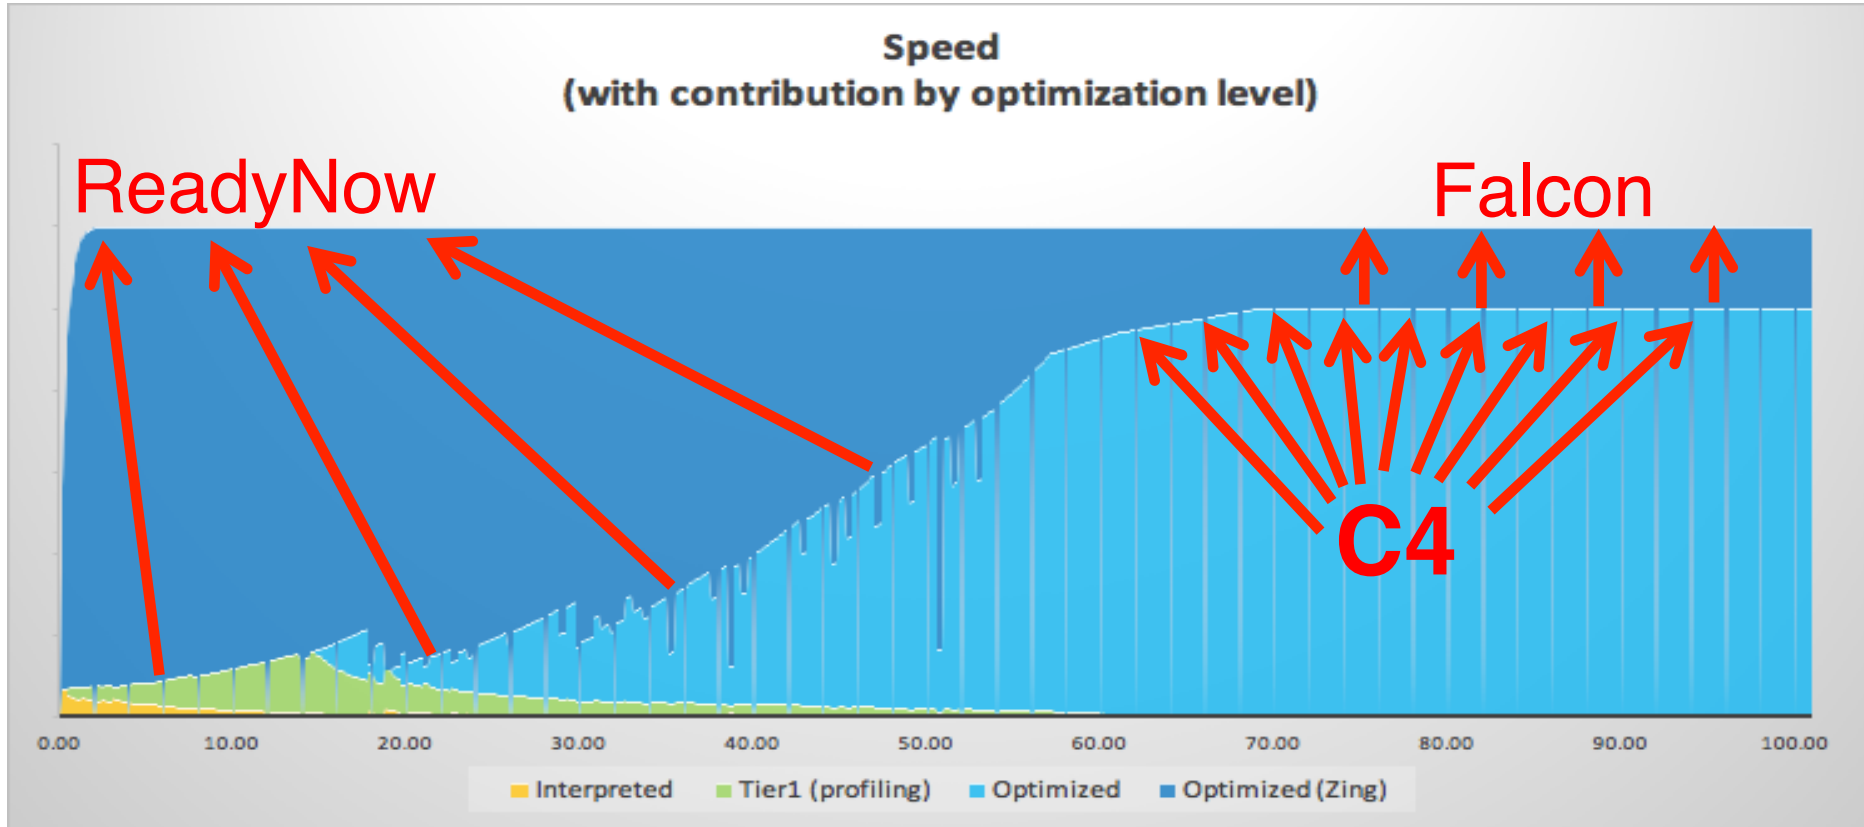
- New and rapidly improving feature for Zing
- Speculative optimisations are responsible for over half of speed improvement in JVM applications
- Zing 'stashes' compiled code for methods
 - Including speculative optimisations
- JVM requests native compiled code for method
 - Provides precise description (profile, types, etc.)



Summary



JVM Performance Graph: Zing



The Zing JVM

- Start fast
- Go faster
- Stay fast
- Simple replacement for other JVMs
 - No recoding necessary

Try Zing free for 30 days:

azul.com/zingtrial

Questions?

© Copyright Azul Systems 2015

Simon Ritter

Deputy CTO, Azul Systems
azul.com



@speakjava

Zing: Supported Platforms

- Linux:
 - RHEL 7.0, 6.0, 5.9 or later
 - CentOS 7.0, 6.0, 5.9 or later
 - Oracle Linux 7.0, 6.0 or later
 - Red Hat MRG Realtime
 - SLES 12 SP1, 11 SP3, SP2 and SP1
 - Ubuntu 16.04 LTS, 14.04 and 12.04 LTS
 - Debian: Wheezy and Jessie
- Hypervisors: VMware, KVM
- Cloud: Amazon AWS, Docker
- Multiple Private Clouds
- Java Versions: 6, 7, 8
- Hardware: Intel/AMD x64

